

# notebook

April 18, 2018

## 1 Domain classification

### 1.0.1 Emanuele Alessi 1486470

```
In [1]: import pickle
import numpy as np
import os
import tqdm
from data_preprocessing import preprocess_text
import seaborn as sn
import pandas as pd
import matplotlib.pyplot as plt

In [2]: embedding = pickle.load(open('vectors.pkl', 'rb'))
emb = {}
row = 0
for word in open('tmp/metadata.tsv', 'r'):
    emb[word.replace('\n', '')] = row
    row += 1

In [3]: def read_data(directory):
    dataset = np.array([])
    stopwords = set([w.rstrip('\r\n') for w in open('stopwords.txt')])
    for domain in os.listdir(directory):
        print(domain)
        files = os.listdir(os.path.join(directory, domain))
        for i in tqdm.tqdm(range(len(files))):
            if files[i].endswith(".txt"):
                with open(os.path.join(directory, domain, files[i]), encoding='utf8') as f:
                    data = []
                    for line in f.readlines():
                        split = preprocess_text(line, stopwords) # split is a list of
                        for sentence in split:
                            if sentence:
                                data += [sentence]
                    val = embeddings_mean(embedding, emb, data, domain)
                    if val is not None:
                        if dataset.size == 0:
```

```

        dataset = np.append(dataset, val)
    else:
        dataset = np.vstack((dataset, val))

    return dataset

def embeddings_mean(embedding, emb, data, domain=None):
    word_count = 0
    s = 0
    for sentence in data:
        for word in sentence:
            if word in emb:
                s += embedding[emb[word]]
            else:
                s += embedding[emb['UNK']]
            word_count += 1
    if word_count > 0:
        if domain:
            return np.append(s / float(word_count), domain)
        else:
            return s / float(word_count)
    return None

def read_test_data(directory):
    dataset = np.array([])
    stopwords = set([w.rstrip('\r\n') for w in open('stopwords.txt')])
    files = os.listdir(directory)
    for i in tqdm.tqdm(range(len(files))):
        if files[i].endswith(".txt"):
            with open(os.path.join(directory, files[i]), encoding='utf8') as file:
                data = []
                for line in file.readlines():
                    split = preprocess_text(line, stopwords) # split is a list of sen
                    for sentence in split:
                        if sentence:
                            data += [sentence]
                val = embeddings_mean(embedding, emb, data)
                if val is not None:
                    if dataset.size == 0:
                        dataset = np.append(dataset, val)
                    else:
                        dataset = np.vstack((dataset, val))

    return dataset

```

```

In [4]: TRAIN_DIR = "dataset/DATA/TRAIN"
        VALID_DIR = "dataset/DATA/DEV"
        TEST_DIR = "dataset/DATA/TEST"
        TMP_DIR = "tmp/"

```

```

if os.path.exists(TMP_DIR + 'train.pkl'):
    # if the train dataset has already been stored in train.pkl then load it
    train = pickle.load(open(TMP_DIR + 'train.pkl', 'rb'))
else:
    # otherwise read the train dataset and store it in train.pkl
    train = read_data(TRAIN_DIR)
    pickle.dump(train, open(TMP_DIR + 'train.pkl', 'wb'))

if os.path.exists(TMP_DIR + 'dev.pkl'):
    # if the validation dataset has already been stored in dev.pkl then load it
    validation = pickle.load(open(TMP_DIR + 'dev.pkl', 'rb'))
else:
    # otherwise read the validation dataset and store it in dev.pkl
    validation = read_data(VALID_DIR)
    pickle.dump(validation, open(TMP_DIR + 'dev.pkl', 'wb'))

```

In [5]: train.shape

Out[5]: (114563, 129)

In [6]: validation.shape

Out[6]: (24546, 129)

```

In [7]: X_train = train[:, :-1].astype(np.float32)
        Y_train = train[:, train.shape[1] - 1]
        del train

        X_validation = validation[:, :-1].astype(np.float32)
        Y_validation = validation[:, validation.shape[1] - 1]
        del validation

```

```

In [8]: from sklearn.model_selection import KFold
        from sklearn.model_selection import cross_val_score
        from sklearn.metrics import classification_report
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import accuracy_score
        from sklearn.linear_model import LogisticRegression
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
        from sklearn.naive_bayes import GaussianNB

```

```

In [9]: # Test to choose the best classification algorithm
        models = []
        models.append(('LR', LogisticRegression()))
        models.append(('LDA', LinearDiscriminantAnalysis()))
        models.append(('KNN', KNeighborsClassifier()))
        models.append(('CART', DecisionTreeClassifier()))

```

```
models.append(('NB', GaussianNB()))
```

```
# evaluation for each model
```

```
for name, model in models:
```

```
    kfold = KFold()
```

```
    cv_results = cross_val_score(model, X_train, Y_train, cv=kfold, scoring='accuracy')
```

```
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
```

```
    print(msg)
```

```
LR: 0.144731 (0.124815)
```

```
LDA: 0.135462 (0.107099)
```

```
KNN: 0.140838 (0.117109)
```

```
CART: 0.101681 (0.081302)
```

```
NB: 0.080645 (0.059856)
```

```
In [10]: # KNN returns the highest accuracy score
```

```
model = KNeighborsClassifier()
```

```
model.fit(X_train, Y_train)
```

```
predictions = model.predict(X_validation)
```

```
array = confusion_matrix(Y_validation, predictions)
```

```
print(classification_report(Y_validation, predictions))
```

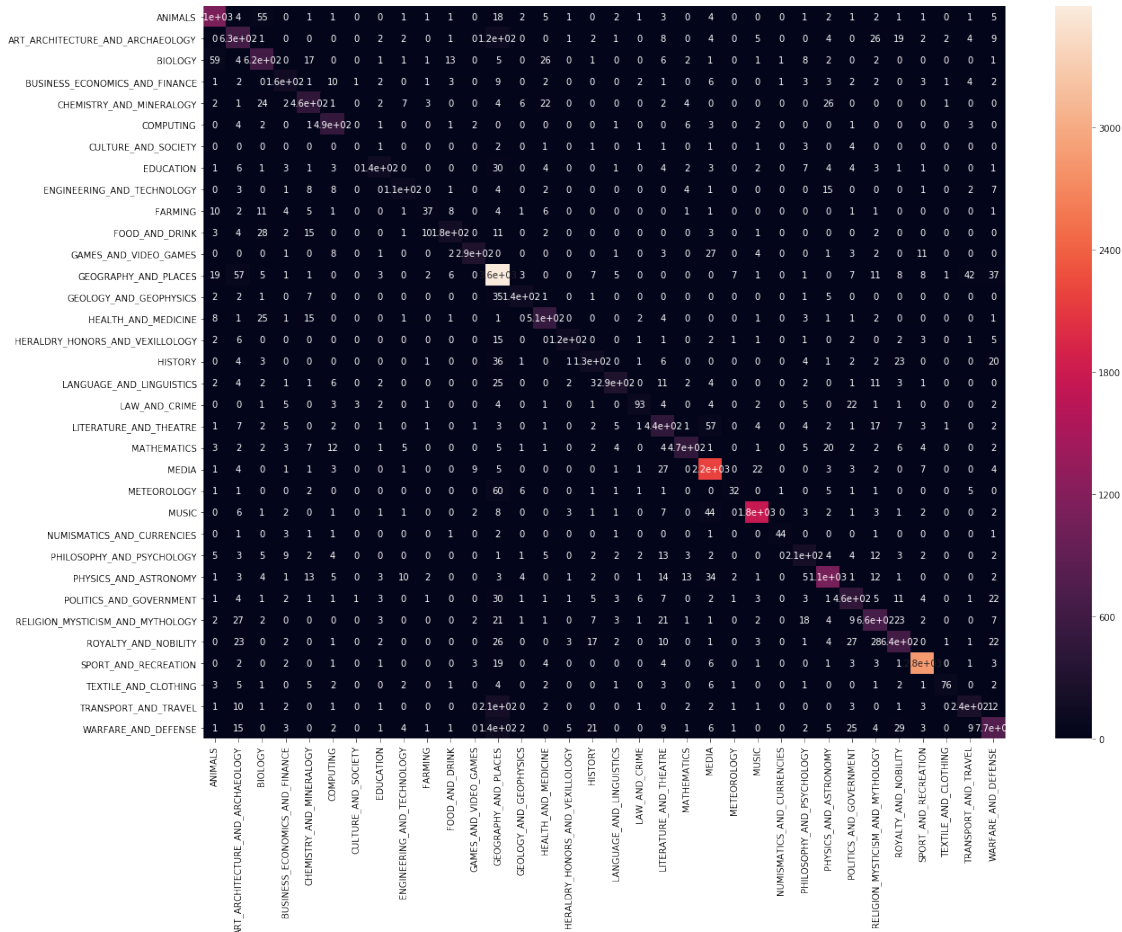
```
print('Accuracy: %.2f' % (accuracy_score(Y_validation, predictions) * 100) + '%')
```

	precision	recall	f1-score	support
ANIMALS	0.90	0.91	0.90	1241
ART_ARCHITECTURE_AND_ARCHAEOLOGY	0.74	0.75	0.75	841
BIOLOGY	0.78	0.80	0.79	776
BUSINESS_ECONOMICS_AND_FINANCE	0.73	0.71	0.72	217
CHEMISTRY_AND_MINERALOGY	0.81	0.81	0.81	569
COMPUTING	0.86	0.95	0.90	515
CULTURE_AND_SOCIETY	0.00	0.00	0.00	16
EDUCATION	0.80	0.63	0.71	222
ENGINEERING_AND_TECHNOLOGY	0.75	0.66	0.70	167
FARMING	0.60	0.39	0.47	95
FOOD_AND_DRINK	0.81	0.68	0.74	258
GAMES_AND_VIDEO_GAMES	0.94	0.82	0.87	354
GEOGRAPHY_AND_PLACES	0.81	0.94	0.87	3827
GEOLOGY_AND_GEOPHYSICS	0.83	0.71	0.77	192
HEALTH_AND_MEDICINE	0.85	0.88	0.87	577
HERALDRY_HONORS_AND_VEXILLOLOGY	0.87	0.74	0.80	166
HISTORY	0.63	0.55	0.58	232
LANGUAGE_AND_LINGUISTICS	0.89	0.78	0.83	376
LAW_AND_CRIME	0.81	0.60	0.69	155
LITERATURE_AND_THEATRE	0.72	0.77	0.74	571
MATHEMATICS	0.92	0.83	0.87	564
MEDIA	0.91	0.96	0.93	2273

METEOROLOGY	0.70	0.27	0.39	119
MUSIC	0.97	0.95	0.96	1861
NUMISMATICS_AND_CURRENCIES	0.94	0.77	0.85	57
PHILOSOPHY_AND_PSYCHOLOGY	0.72	0.71	0.71	294
PHYSICS_AND_ASTRONOMY	0.90	0.89	0.90	1223
POLITICS_AND_GOVERNMENT	0.78	0.79	0.78	580
RELIGION_MYSTICISM_AND_MYTHOLOGY	0.81	0.81	0.81	823
ROYALTY_AND_NOBILITY	0.82	0.79	0.80	811
SPORT_AND_RECREATION	0.98	0.98	0.98	2899
TEXTILE_AND_CLOTHING	0.92	0.64	0.75	119
TRANSPORT_AND_TRAVEL	0.77	0.49	0.60	495
WARFARE_AND_DEFENSE	0.82	0.73	0.77	1061
avg / total	0.86	0.86	0.85	24546

Accuracy: 85.72%

```
In [11]: conf_mat = pd.DataFrame(array,
                                index = [domain for domain in os.listdir(TRAIN_DIR)],
                                columns = [domain for domain in os.listdir(TRAIN_DIR)])
plt.figure(figsize = (20, 15))
sn.heatmap(conf_mat, annot=True)
plt.savefig('conf_mat.png')
```



```
In [12]: if os.path.exists(TMP_DIR + 'test.pkl'):
# if the test dataset has already been stored in test.pkl then load it
test = pickle.load(open(TMP_DIR + 'test.pkl', 'rb'))
else:
# otherwise read the test dataset and store it in test.pkl
test = read_test_data(TEST_DIR)
pickle.dump(train, open(TMP_DIR + 'test.pkl', 'wb'))

In [13]: test.shape

Out[13]: (24526, 128)

In [14]: test_predictions = model.predict(test)

In [15]: with open('test_answers.tsv', 'w') as f:
for i in range(len(test_predictions)):
f.write('test_' + str(i) + '\t' + test_predictions[i] + '\n')
```